

HIGH PERFORMANCE COMPUTING

Assignment

Do data access analysis for project problem.

Propose optimal data access technique for your project problem.

Provide a detailed analysis report.

Project problem - Using Discrete Fourier Transform for finding the normalized frequency of a random wave and parallelizing the code using Openmp.

Bhaavanaa Thumu - CED17I021

Machine balance

The machine balance of a processor chip is the ratio of possible memory bandwidth in GWords/sec to peak performance in GFlops/sec.

$$\text{Memory bandwidth} = 34.1 / 8 = 4.2625$$

$$\text{No. of cores} = 2$$

$$\text{Processor speed} = 2.7 \text{ GHz}$$

$$\text{Flops} = 4$$

$$\text{Peak performance} = \text{no. of cores} \times \text{processor speed} \times \text{flops}$$

$$= 2 \times 2.7 \times 4 = 21.6$$

$$\text{Machine balance} = \text{memory bandwidth} / \text{peak performance}$$

$$= 4.2625 / 21.6$$

$$= 0.1973$$

Hence, machine balance is 0.1973.

Code balance

The code balance of a loop is the ratio of data traffic to the number of floating point operations.

$$\text{No. of load operations in the program} = 4$$

$$\text{No. of store operation in the program} = 7$$

$$\text{Total number of memory accesses} = 4 + 7 = 11$$

$$\text{No. of floating point operations} = 18$$

Code balance = total no. of memory accesses / no. of floating point operations

$$= 11 / 18$$

$$= 0.6111$$

Hence, code balance is 0.6111.

Lightspeed of loop

$$\text{machine balance} / \text{code balance} = 0.1973 / 0.6111 = 0.3228$$

$$\text{Lightspeed of loop} = \min(1, \text{machine balance/code balance})$$

$$= \min(1, 0.3228)$$

$$= 0.3228$$

Hence, the lightspeed of loop is 0.3228.

Algorithm classification and access optimizations

The optimization potential of many loops on cache-based processors can easily be estimated just by looking at basic parameters like the scaling behavior of data transfers and arithmetic operations versus problem size.

$O(N) / O(N)$

If both the number of arithmetic operations and the number of data transfers (loads/stores) are proportional to the problem size (or “loop length”) N , optimization potential is usually very limited. Scalar products, vector additions, and sparse matrix-vector multiplication are examples for this kind of problems. They are inevitably memory-bound for large N , and compiler-generated code achieves good performance because $O(N)/O(N)$ loops tend to be quite simple and the correct software pipelining strategy is obvious. Loop nests, however, are a different matter. But even if loops are not nested there is sometimes room for improvement. As an example, consider the following from the code:

```
for (int k = 0; k < N; k++) // finding the dft
{
    for (int n = 0; n < N; n++) // dft value
    {
        dft_val[k].real += samples[n] * cos((2 * PI * k * n) / N);
        dft_val[k].img += samples[n] * sin((2 * PI * k * n) / N);
    }

    k_by_N[k] = (double)k/N; // k/N value

    power_spectrum[k] = dft_val[k].real*dft_val[k].real +
dft_val[k].img*dft_val[k].img; // power spectrum value

    sum_power_spectrum = sum_power_spectrum +
power_spectrum[k]; // sum of power spectrum values
    sum_freq_power = sum_freq_power +
2*PI*k_by_N[k]*power_spectrum[k]; // sum of freq*power_spectrum
values
}
```

The above code removes the need for loading `dft_val[k].real` and `dft_val[k].img`, as it is used for computing as soon as it is updated. We are fusing the loops to reduce the loads. Similarly, for producing the arrays, `sum_power_spectrum` and `sum_freq_power_spectrum`, also use the recently used data for their computation. This gives a room for improvement.