

# HIGH PERFORMANCE COMPUTING PRACTICE

---

## EXPERIMENT - 8

### CUDA:

- 1) Matrix addition
- 2) Matrix multiplication column major order
- 3) Matrix multiplication block based approach

Use input as a larger double number (64-bit).

Run experiment for Threads = {1,2,4,8,16,32,64,128,256,500}.

Estimate the parallelization fraction.

Document the report and submit.

**Bhaavanaa Thumu - CED17I021**

---

---

## Specifications

The size of all the matrices used for below programs is 1,000x1,000.

All the elements of the matrices are produced randomly and are large double numbers (64-bit).

## Formulae used

Speed-up =  $T(1)/T(P)$

Parallel fraction,  $f = (1 - T(P)/T(1))/(1 - (1/P))$

where,

$T(1)$  - time taken for serial execution

$T(P)$  - time taken for parallel execution

$P$  - number of processes/threads/processors

## Other functions used

cudaMalloc - allocates a block of memory

cudaMemcpy - copies 'n' characters from source to destination memory area.

---

## 1. Matrix addition


This program contains matrices of size 1,000x1,000.

The number of processors initialized and the execution time in each case is as shown in the table.

<i>No. of processors</i>	<i>Execution time</i>	<i>Speed-up</i>	<i>Parallel fraction</i>
1	0.175272	1	#DIV/0!
2	0.169077	1.036640111	0.07069012734
4	0.172974	1.013285234	0.01748140034
8	0.169208	1.035837549	0.03954017592
16	0.177450	0.98772612	-0.01325482678
32	0.170749	1.026489174	0.02663804387
64	0.169748	1.032542357	0.03201699381
128	0.173706	1.009015233	0.009005036222
256	0.171361	1.022823163	0.02240139472
500	0.169900	1.031618599	0.03071092662

---

The plot below is the graph between no. of processors and execution time.



x-axis : no. of processors

y-axis : execution time

Hence, in this case, the optimal number of processors is 2. In this case-

Execution time = 0.169077

Speed-up = 1.036640111

Parallel fraction = 0.07069012734

---

## **Matrix addition code**

```
%%cu

#include<stdio.h>
#include<math.h>
#include <sys/time.h>

#define M 1000
#define N 1000

__global__ void matrices_add(double *d_result, double *d_a,
double *d_b, int thread_count)
{
    int elementsReceived = M/thread_count;

    if(thread_count != 1)
    {
        int start_index, end_index;

        start_index = (threadIdx.x * elementsReceived);

        if(threadIdx.x == thread_count-1)
            end_index = M;
        else
            end_index=(threadIdx.x + 1)*elementsReceived;

        for(int i=start_index; i<end_index; i++)
        {
            for(int j=0; j<N; j++)
            {
                d_result[i*N+j] = d_a[i*N+j] + d_b[i*N+j];
            }
        }
    }
}
```

---

```

        }
    }
    else
    {
        for(int i=0; i<M; i++)
        {
            for(int j=0; j<N; j++)
                d_result[i*N+j] = d_a[i*N+j] + d_b[i*N+j];
        }
    }
}

int main()
{
    double *a, *b, *out, *d_a, *d_b, *d_result;
    struct timeval t1, t2;

    int size = sizeof(double)*M*N;
    a = (double*)malloc(size);
    b = (double*)malloc(size);
    out = (double*)malloc(size);

    cudaMalloc((void**)&d_a, size);
    cudaMalloc((void**)&d_b, size);
    cudaMalloc((void**)&d_result, size);
    gettimeofday(&t1, 0);

    for(int i = 0; i < M; i++)
    {
        for(int j=0;j<N;j++)
        {
            a[i*N+j] = pow(2,15)+rand()+0.13246549884;
            b[i*N+j] = pow(2,15)+rand()+0.13246549884;

```

---

---

```
        // a[i*N+j] = i;
        // b[i*N+j] = j;
    }
}

cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);

int thread_count = 500;

matrices_add<<<1,thread_count>>>(d_result, d_a, d_b,
thread_count);

cudaMemcpy(out, d_result, size, cudaMemcpyDeviceToHost);

gettimeofday(&t2, 0);
double time = (float)(1000000.0*(t2.tv_sec-t1.tv_sec) +
t2.tv_usec-t1.tv_usec)/1000000.0;
printf("time taken : %lf sec\n", time);
return 0;
}
```

---

## 2. Matrix multiplication column major order

This program contains matrices of size 1,000x1,000.


The number of processors initialized and the execution time in each case is as shown in the table.

<i>No. of processors</i>	<i>Execution time</i>	<i>Speed-up</i>	<i>Parallel fraction</i>
1	0.000947	1	#DIV/0!
2	0.000950	0.9968421053	-0.006335797254
4	0.000956	0.9905857741	-0.01267159451
8	0.000934	1.01391863	0.01568864082
16	0.000972	0.9742798354	-0.02815909891
32	0.000963	0.9833852544	-0.01744047416
64	0.000939	1.008519702	0.00858182062
128	0.000946	1.001057082	0.001064280904
256	0.000953	0.9937040923	-0.006360643518
500	0.000950	0.9968421053	-0.003174247121



---

The plot below is the graph between no. of processors and execution time.



x-axis : no. of processors

y-axis : execution time

Hence, in this case, the optimal number of processors is 8. In this case-

Execution time = 0.000934

Speed-up = 1.01391863

Parallel fraction = 0.01568864082

---

## Matrix multiplication column major order code

```
%%cu

#include<stdio.h>
#include<math.h>
#include <sys/time.h>

#define M 1000
#define N 1000

#define thread_count 500

__global__ void mat_mul(double *d_result, double *d_a, double
*d_b)
{
    int elementsReceived = M/thread_count;
    if(thread_count != 1)
    {
        int start_index, end_index;

        start_index=(threadIdx.x * elementsReceived);

        if(threadIdx.x==thread_count-1)
            end_index = M;
        else
            end_index=(threadIdx.x + 1)*elementsReceived;

        for(int i=start_index; i<end_index; i++)
        {
            for(int j=0; j<N; j++)
            {
```

---

```

        for(int k=0; k<N; k++)
        {
            d_result[i+j*M] = d_result[i+j*M] +
d_a[i+k*M] * d_b[k+j*N];    // Take row of a and column of b
        }
    }
}
else
{
    for(int i=0; i<M; i++)
    {
        for(int j=0; j<N; j++)
        {
            for(int k=0; k<N; k++)
            {
                d_result[i+j*M] = d_result[i+j*M] +
d_a[i+k*M] * d_b[k+j*N];    //Take row of a and column of b
            }
        }
    }
}
}

int main(void)
{
    double *a, *b, *out, *d_a, *d_b, *d_result;
    struct timeval t1, t2;
    int size = sizeof(double)*M*N;
    a = (double *)malloc(size);
    b = (double *)malloc(size);
    out = (double *)malloc(size);

```

---

---

```
    cudaMalloc((void **)&d_a, size);
    cudaMalloc((void **)&d_b, size);
    cudaMalloc((void **)&d_result, size);

    for(int i = 0; i < M; i++)
    {
        for(int j=0;j<N;j++)
        {
            a[i+j*M] = pow(2,15)+rand()+0.13246549884;
            b[i+j*M] = pow(2,15)+rand()+0.13246549884;
            // a[i+j*M] = i;
            // b[i+j*M] = j;
        }
    }

    gettimeofday(&t1, 0);
    cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);
    mat_mul<<<1,thread_count>>>(d_result, d_a, d_b);

    cudaMemcpy(out, d_result, size, cudaMemcpyDeviceToHost);
    gettimeofday(&t2, 0);
    double time = (float)(1000000.0*(t2.tv_sec-t1.tv_sec) +
t2.tv_usec-t1.tv_usec)/1000000.0;
    printf("time taken : %lf sec\n", time);
    return 0;
}
```

---

### 3. Matrix multiplication block based approach


This program contains matrices of size 1,000x1,000.

The number of processors initialized and the execution time in each case is as shown in the table.

<i>No. of processors</i>	<i>Execution time</i>	<i>Speed-up</i>	<i>Parallel fraction</i>
1	0.000962	1	#DIV/0!
2	0.000969	0.9927760578	-0.01455301455
4	0.000966	0.9958592133	-0.005544005544
8	0.000988	0.9736842105	-0.03088803089
16	0.000956	1.006276151	0.006652806653
32	0.000958	1.004175365	0.004292133324
64	0.000952	1.010504202	0.01056001056
128	0.000968	0.9938016529	-0.006286116522
256	0.000970	0.9917525773	-0.008348620113
500	0.000963	0.9989615784	-0.001041584208

---

The plot below is the graph between no. of processors and execution time.



x-axis : no. of processors

y-axis : execution time

Hence, in this case, the optimal number of processors is 64. In this case-

Execution time = 0.000952

Speed-up = 1.010504202

Parallel fraction = 0.01056001056

---

## **Matrix multiplication block based approach code**

```
%%cu

#include <stdio.h>
#include<sys/time.h>

#define M 1000
#define N 1000

#define num_threads 500

__global__ void matmul(double *a, double *b, double *c)
{
    int B=5;
    int elementsReceived = N/num_threads;
    int mini,minik;
    double r;

    if(num_threads != 1)
    {
        int start_index,end_index;
        start_index=(threadIdx.x * elementsReceived);

        if(threadIdx.x==num_threads-1)
            end_index=N;//Num of rows
        else
            end_index=(threadIdx.x + 1)*elementsReceived;

        for(int jj=0; jj<N; jj=jj+B)
        {
```

---

```

    for(int kk=0; kk<N; kk=kk+B)
    {
        for(int i=start_index; i<end_index; i++)
        {
            if (jj+B<N)
                mini=jj+B;
            else
                mini=N;

            for(int j=jj; j<mini; j++)
            {
                r=0;
                if (kk+B<N)
                    minik=kk+B;
                else
                    minik=N;

                for(int k=kk; k<minik; k++)
                    r+=a[i*N+k]* b[k*N+j];

                c[i*N+j]+=r;
            }
        }
    }
}
else
{
    for(int jj=0; jj<N; jj=jj+B)
    for(int kk=0; kk<N; kk=kk+B)
    for(int i=0; i<N; i++)
    {
        if (jj+B<N)

```

---



---

```

        mini=jj+B;
    else
        mini=N;
    for(int j=jj;j<mini;j++)
    {
        r=0;
        if (kk+B<N)
            minik=kk+B;
        else
            minik=N;

        for(int k=kk;k<minik;k++)
            r+=a[i*N+k]* b[k*N+j];

        c[i*N+j]+=r;
    }
}

}

int main(void)
{
    double *a, *b, *c;    // host copies of a, b, c
    double *d_a, *d_b, *d_c; // device copies of a, b, c
    int size = sizeof(double)*N*N;
    double start,end;
    cudaMalloc((void **)&d_a, size);    // Allocate space for
device copies of a, b, c
    cudaMalloc((void **)&d_b, size);
    cudaMalloc((void **)&d_c, size);

    a = (double *)malloc(size);
    b = (double *)malloc(size);

```

---

---

```
c = (double *)malloc(size);

for(int i = 0; i < N; i++)           // Setup input values
{
    for(int j=0;j<N;j++)
    {
        a[i+j*N] = pow(2,15)+rand()+0.13246549884;
        b[i+j*N] = pow(2,15)+rand()+0.13246549884;
    }
}

cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);    // Copy
inputs to device
cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);

matmul<<<1,num_threads>>>(d_a, d_b, d_c);           // Launch
matmul() kernel on GPU

cudaMemcpy(c, d_c, size, cudaMemcpyDeviceToHost);

return 0;
}
```